



ÉCOLE POLYTECHNIQUE
FÉDÉRALE DE LAUSANNE

What is Reactive Programming?

Principles of Reactive Programming

Martin Odersky

Changing Requirements...

	<i>10 years ago</i>	<i>Now</i>
Server nodes	10's	1000's
Response times	seconds	milliseconds
Maintenance downtimes	hours	none
Data volume	GBs	TBs → PBs.

... Need New Architectures

Previously: Managed servers and containers.

Now: Reactive applications

- ▶ event-driven
- ▶ scalable
- ▶ resilient
- ▶ responsive

Reactive

[Merriam Webster] reactive: *“readily responsive to a stimulus”*.

- ▶ React to events (*event-driven*)
- ▶ React to load (*scalable*)
- ▶ React to failures (*resilient*)
- ▶ React to users (*responsive*)

Event-Driven

Traditionally: Systems are composed of multiple threads, which communicate with shared, synchronized state.

- ▶ Strong coupling, hard to compose.

Now: Systems are composed from loosely coupled event handlers.

- ▶ Events can be handled asynchronously, without blocking.

Scalable

An application is *scalable* if it is able to be expanded according to its usage.

- ▶ scale up: make use of parallelism in multi-core systems
- ▶ scale out: make use of multiple server nodes

Important for scalability: Minimize shared mutable state.

Important for scale out: Location transparency, resilience.

Resilient

An application is *resilient* if it can recover quickly from failures.

Failures can be:

- ▶ software failures
- ▶ hardware failures, or
- ▶ connection failures.

Typically, resilience cannot be added as an afterthought; it needs to be part of the design from the beginning.

Needed:

- ▶ loose coupling.
- ▶ strong encapsulation of state.
- ▶ pervasive supervisor hierarchies.

Responsive

An application is *responsive* if it provides rich, real-time interaction with its users even under load and in the presence of failures.

Responsive applications can be built on an event-driven, scalable, and resilient architecture.

Still need careful attention to algorithms, system design, back-pressure, and many other details.

Call-backs

Handling events is often done using call-backs. E.g. using Java observers:

```
class Counter implements ActionListener {  
    private var count = 0  
    button.addActionListener(this)  
  
    def actionPerformed(e:(ActionEvent)): Unit = {  
        count += 1  
    }  
}
```

Call-backs

Handling events is often done using call-backs. E.g. using Java observers:

```
class Counter implements ActionListener {  
    private var count = 0  
    button.addActionListener(this)  
  
    def actionPerformed(e: ActionEvent): Unit = {  
        count += 1  
    }  
}
```

Problems:

- ▶ needs shared mutable state.
- ▶ cannot be composed.
- ▶ leads quickly to “call-back hell”.

How To Do Better

Use fundamental constructions from functional programming ...

How To Do Better

Use fundamental constructions from functional programming ...

... to get *composable* event abstractions.

How To Do Better

Use fundamental constructions from functional programming ...
... to get *composable* event abstractions.

- ▶ Events are first class.
- ▶ Events are often represented as messages.
- ▶ Handlers of events are also first-class.
- ▶ Complex handlers can be composed from primitive ones.

Contents of This Course

- ▶ Review of functional programming
- ▶ An important class of functional patterns: *monads*
- ▶ Functional programs in a stateful world
- ▶ Abstracting over events: *futures*
- ▶ Abstracting over event streams: *observables*
- ▶ Message passing architecture: *actors*
- ▶ Handling failures: *supervisors*
- ▶ Scaling out: *distributed actors*

Prerequisites

- ▶ Need a solid grounding in functional programming.
- ▶ Ideally, the “Principles of Functional Programming in Scala” class.
- ▶ If you know some other functional language, the switch should be easy.